

Predicting Customer Churn

This mini-project analyzes and predicts **Customer Churn** using the [Credit Card Customers](#) dataset from [Kaggle](#).

The code is in the repository [customer_churn_production](#), specifically in the notebook [churn_notebook.ipynb](#).

The **main goal has been to find a model which best predicts customer churn**; model interpretation has been considered as a secondary objective.

In order to achieve that, the typical steps from the data science process have been followed:

- Dataset loading
- Exploratory Data Analysis (EDA) and Data Cleaning
- Feature Engineering (FE)
- Model Training
- Generation of classification report plots (evaluation) as well as feature importance plots (interpretation).

Dataset

The dataset consists in a data frame with 10127 rows (data points) and 22 columns (features for each data point).

The feature columns are the following:

- **Target** (renamed to `Churn = {0,1}`)
 - `Attrition_Flag: Existing Customer, Attrited Customer` (churn).
- **Dropped** (irrelevant information)
 - `Unnamed: 0`: copy of index (not present in the current Kaggle dataset).
 - `CLIENTNUM`: client ID.
- **Categorical** (encoded as churn ratios per category level)
 - `Gender`: F, M.
 - `Education_Level`: High School, Graduate, etc.
 - `Marital_Status`: Married, Single, Divorced, Unknown
 - `Income_Category`: annual income range: < \$40K, \$40K – 60K, \$60K – \$80K, \$80K–\$120K, > 120k
 - `Card_Category`: type of card/product: Blue, Silver, Gold, Platinum
- **Numerical** (transformed if the absolute skew is larger than 0.75)
 - `Customer_Age`: customer's age in years (quite normal).
 - `Dependent_count`: number of dependents (quite normal).
 - `Months_on_book`: period of relationship with bank (normal with peak).
 - `Total_Relationship_Count`: total number of products held by the customer (quite uniform).
 - `Months_Inactive_12_mon`: number of months inactive in the last 12 months (quite normal).
 - `Contacts_Count_12_mon`: number of Contacts in the last 12 months (quite normal).

- **Credit_Limit**: credit limit on the credit card (exponential decaying).
- **Total_Revolving_Bal**: total revolving balance on the credit card, i.e., amount that goes unpaid at the end of the billing cycle (non-normal; peaks at both tails).
- **Avg_Open_To_Buy**: open to buy credit line (average of last 12 months) (exponential).
- **Total_Amt_Chng_Q4_Q1**: change in transaction amount (Q4 over Q1) (quite normal, maybe large central peak).
- **Total_Trans_Amt**: total transaction amount (last 12 months) (tri-modal).
- **Total_Trans_Ct**: total transaction count (last 12 months) (bi-modal).
- **Total_Ct_Chng_Q4_Q1**: change in transaction count (Q4 over Q1) (quite normal, maybe large central peak).
- **Avg_Utilization_Ratio**: average card utilization ratio (exponential decaying).

Data Processing

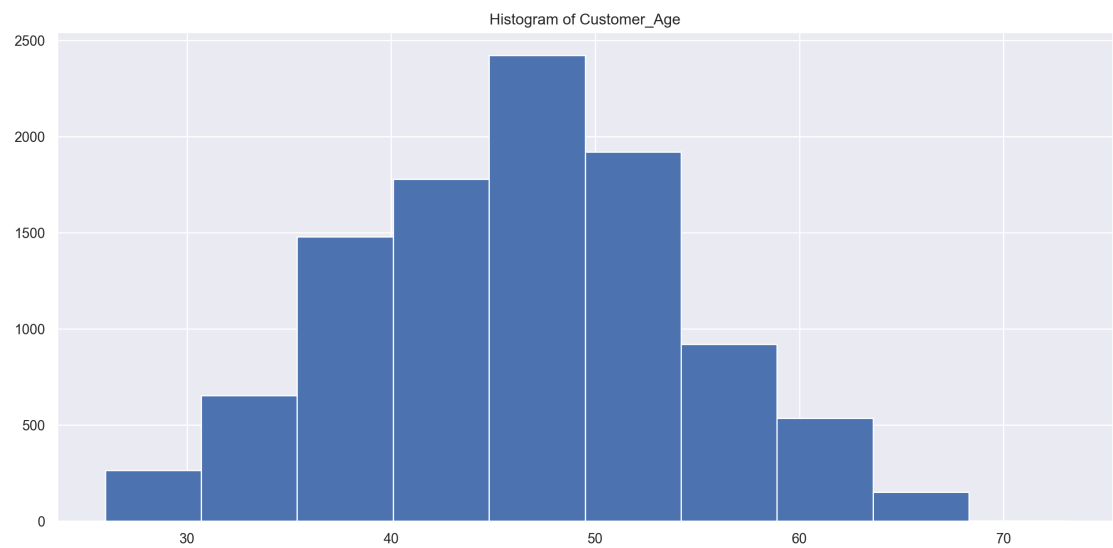
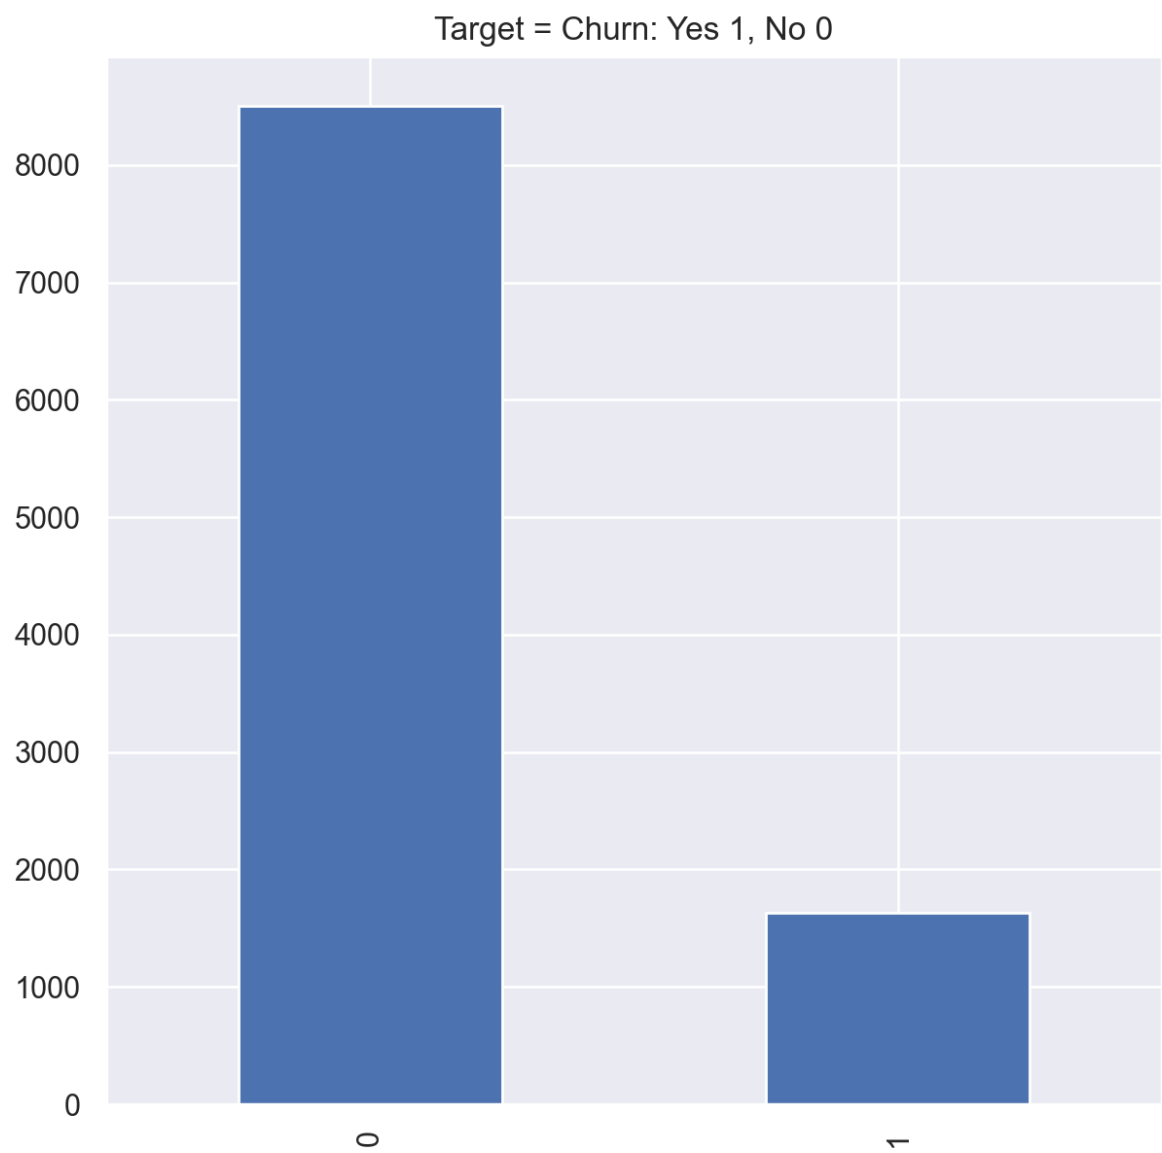
First, some initial data inspection and **data cleaning** have been carried out:


- Duplicates have been checked (no duplicates found).
- Missing data have been checked (no missing data).
- Unnecessary columns have been identified and dropped.
- Target variable has been mapped to integer labels.

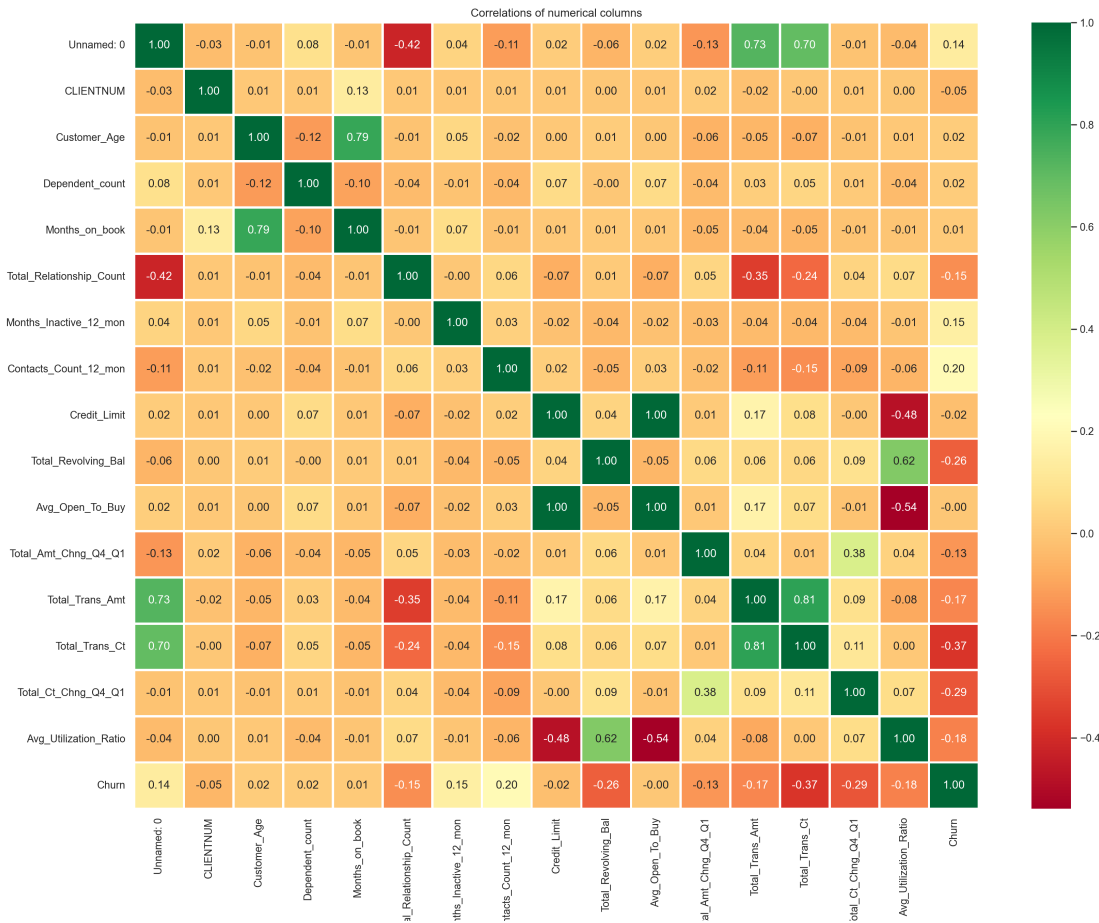
After that, the categorical and numerical columns have been automatically detected and the most interesting ones plotted to build an **Exploratory Data Analysis**:

- Distributions for numerical values
- Correlation heatmap for numerical values
- Bar plot for categorical values

In the following, some plots are provided:







Finally, **feature engineering** has been performed:

- All numerical features have been checked for their skewness; if the absolute value was larger than 0.75, they have been transformed with a Yeo-Johnson transformation (parameters learned and preserved).
- All categorical values have been encoded as churn ratios: for each category level in a categorical feature, the churn ratio of the associated level has been computed and the category label replaced with that ratio. Another option would have been to introduce one-hot encoded variables.

Before feeding the data to the models, polynomial features and scaling have been added; however, that has been done inside a pipeline optimized with a grid search, in order to obtain the best polynomial degree for the goal: best predictive model.

Data Modeling

After the data exploration, a train/test split has been performed with stratification (related to the target value) and three classifiers have been trained using a grid search with cross validation for finding the best set of hyperparameters:

- Logistic regression as a baseline. 24 variations have been optimized, modifying the regularization strength, the regularization penalty type and the polynomial degree.
- A Support Vector Machine model. 32 variations have been optimized, modifying the regularization strength, the kernel, the kernel factor gamma and the polynomial degree.

- A Random Forest model. 24 variations have been optimized, modifying the information gain criterion, the maximal depth of each tree, the maximal features of each tree and the number of estimators/trees.

The best model parameters have been chosen so that we get the best ROC AUC score on the training split with 5-fold cross validation; the final results are the following:

Logistic Regression:

- Best params: {'model__C': 10.0, 'model__penalty': 'l1', 'polynomial_features__degree': 2}
- Best ROC AUC (train split with CV): 0.9667898323843417

Support Vector Machine:

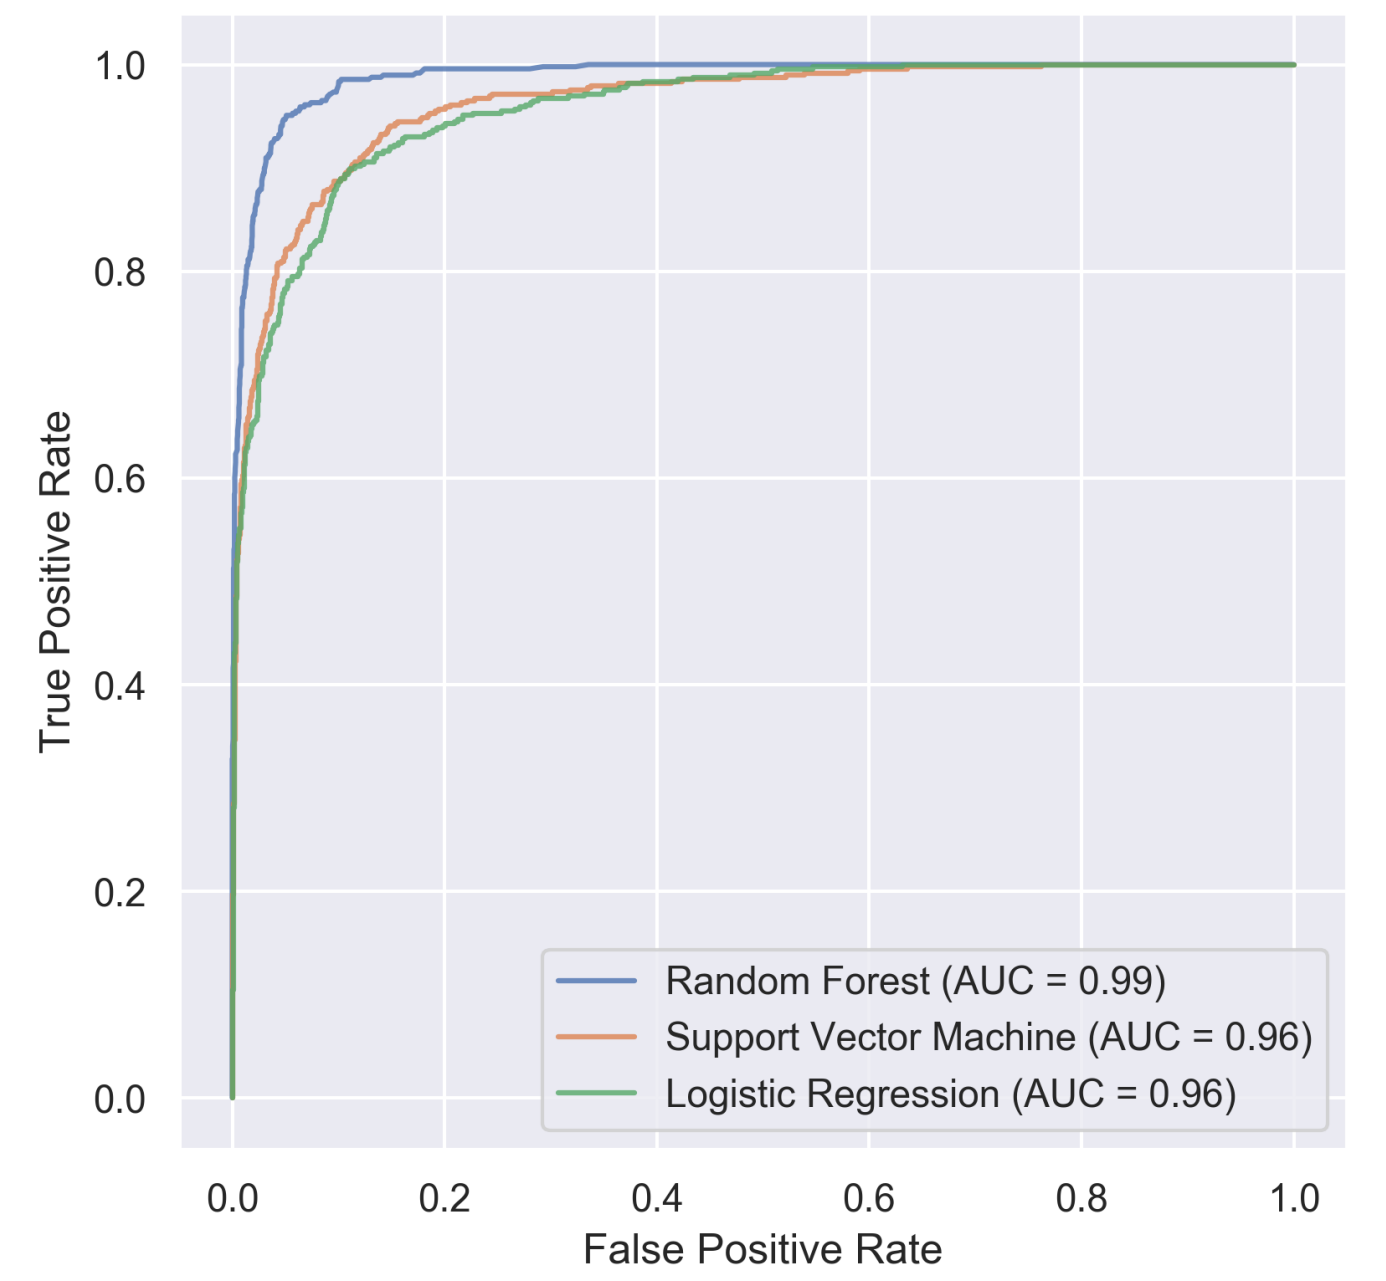
- Best params: {'model__C': 10.0, 'model__gamma': 'auto', 'model__kernel': 'rbf', 'polynomial_features__degree': 1}
- Best ROC AUC (train split with CV): 0.9679994822827043

Random Forest:

- Best params: {'model__criterion': 'entropy', 'model__max_depth': 100, 'model__max_features': 'auto', 'model__n_estimators': 500}
- Best ROC AUC (train split with CV): 0.9895865534929138

Key Findings

The random forest model outperforms the other two models in both the train and test splits when the ROC AUC score is observed; in the following, the ROC curve of all models is shown:



The classification reports from Scikit-Learn reveal also better precision and recall metrics for the random forest model:

--- logistic regression results ---				
TEST results				
	precision	recall	f1-score	support
0	0.95	0.97	0.96	2551
1	0.82	0.72	0.76	488
accuracy			0.93	3039
macro avg	0.88	0.84	0.86	3039
weighted avg	0.93	0.93	0.93	3039
TRAIN results				

	precision	recall	f1-score	support
0	0.96	0.98	0.97	5949
1	0.86	0.79	0.83	1139
accuracy			0.95	7088
macro avg	0.91	0.88	0.90	7088
weighted avg	0.95	0.95	0.95	7088

--- random forest results ---

TEST results

	precision	recall	f1-score	support
0	0.96	0.99	0.97	2551
1	0.92	0.81	0.86	488
accuracy			0.96	3039
macro avg	0.94	0.90	0.92	3039
weighted avg	0.96	0.96	0.96	3039

TRAIN results

	precision	recall	f1-score	support
0	1.00	1.00	1.00	5949
1	1.00	1.00	1.00	1139
accuracy			1.00	7088
macro avg	1.00	1.00	1.00	7088
weighted avg	1.00	1.00	1.00	7088

--- support vector machines results ---

TEST results


	precision	recall	f1-score	support
0	0.95	0.97	0.96	2551
1	0.82	0.74	0.78	488
accuracy			0.93	3039
macro avg	0.89	0.86	0.87	3039
weighted avg	0.93	0.93	0.93	3039

TRAIN results

	precision	recall	f1-score	support
0	0.99	0.99	0.99	5949
1	0.97	0.94	0.95	1139

accuracy			0.99	7088
macro avg	0.98	0.97	0.97	7088
weighted avg	0.99	0.99	0.99	7088

The following diagram shows the distribution of importance values associated with each feature that determine whether a customer is going to churn:



Conclusions and Outlook

In this mini-project, I researched the [Credit Card Customers](#) dataset from [Kaggle](#) with the goal of building a model able to predict **Customer Churn**.

The following steps have been carried out:

- Dataset loading
- Exploratory Data Analysis (EDA) and Data Cleaning
- Feature Engineering (FE)
- Model Training: Logistic Regression, Support Vector Machines, Random Forests.
- Generation of classification report plots (evaluation) as well as feature importance plots (interpretation).

The final best model is a random forest, and it achieves with the test split:

- Precision: 0.94
- Recall: 0.90
- ROC-AUC: 0.99

The classifier seems good enough to be deployed. However, future work involves exploring the following action items to find out whether they bring improvements:

- Add more variables and levels to the grid search.
- Consider using polynomial features in random forests.
- Apply feature selection with LASSO and check its effect.
- Try other classifiers, such as gradient boosting.